Figure 5 The pathology belongs to Config A, not to the missing normalization

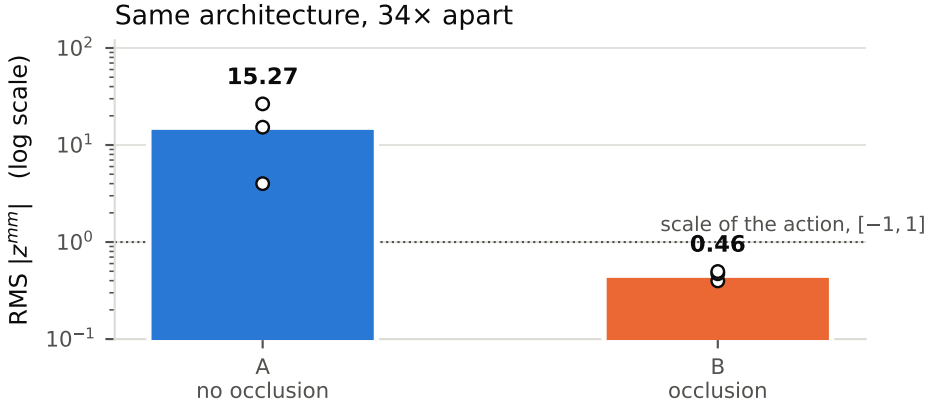

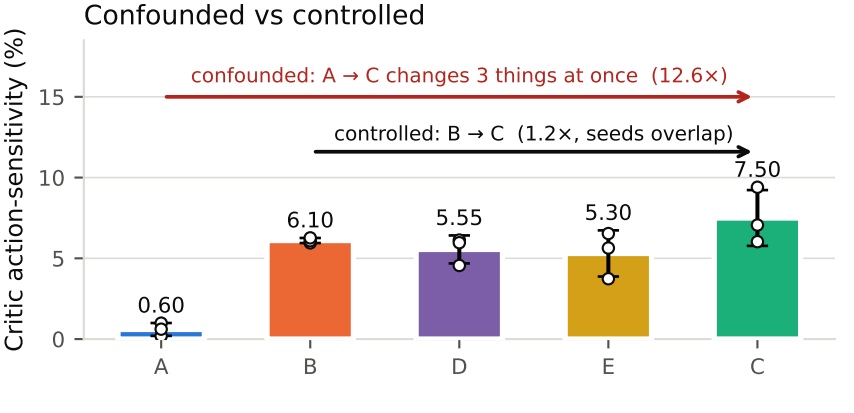

A and B are the same network — `concat`, `normalization: none` — differing only in the occluder. A's fused vector is unbounded (RMS 15.27 ± 11.27, seeds at 4.0 / 15.3 / 26.5, max 171) and 0.60 % action-sensitive; B's is well-scaled (0.46 ± 0.05, max 2.8) and already 6.10 % sensitive — as sensitive as the full method. The Oct 29 audit's "50× improvement from normalization" compared A against C and so changed occlusion, normalization and dualization together.